

# Fuzzing Smart Contract Systems Effectively

Valentin Wüstholtz



# Wish list for fuzzing contracts

1. Find bugs (that is, property violations)
2. Achieve high code coverage
3. Write as little test code as possible
4. Do so quickly :)

**Good news:** there are fuzzers to help you!

**Bad news:** mastering these can take a bit of time

# Overview

1. The fuzzer spectrum
2. Fuzzing contract systems
3. How to increase coverage using fuzzing lessons
4. Testing your specifications and fuzzing setup
5. Fuzzing under the hood

# The fuzzer spectrum

Fuzzing functions vs. fuzzing contract systems

TEST CODE

↑ TEST SUIT

● FOUNDRY/  
ECHIDNA

PROCESS A

- 1) WRITE TEST FOR FUNCTION F
- 2) MAKE SOME INPUTS FUZZABLE
- 3) RUN FUZZER ON F

DF (HARVEY)

→ TEST TIME

PROCESS B

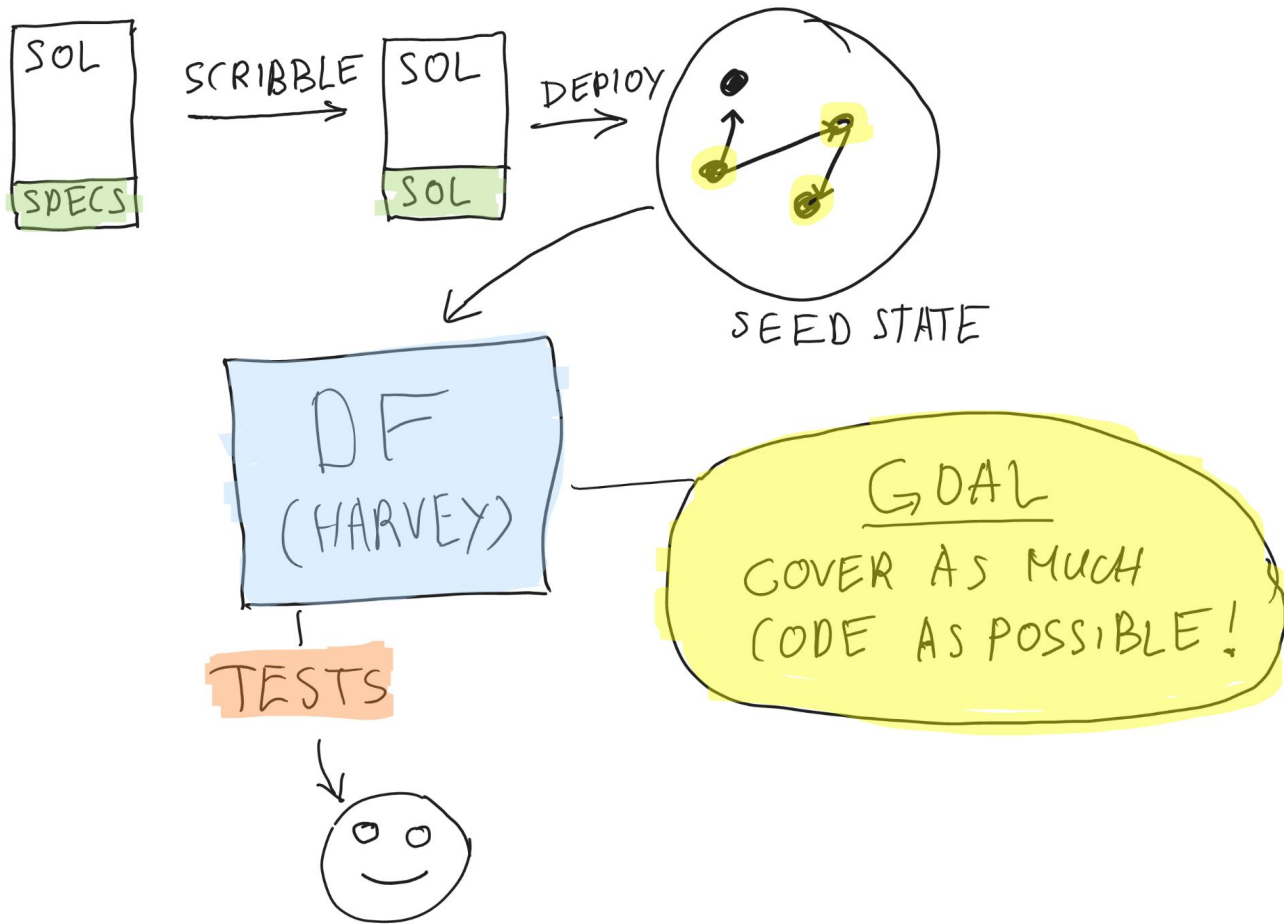
- 1) DEPLOY CONTRACT SYSTEM S
- 2) RUN FUZZER ON S

# What makes fuzzing systems hard?

- Transactions can invoke **many contracts**
- Transactions can invoke **many functions**
- Functions can have **many inputs**
- **Grows exponentially** with number of transactions

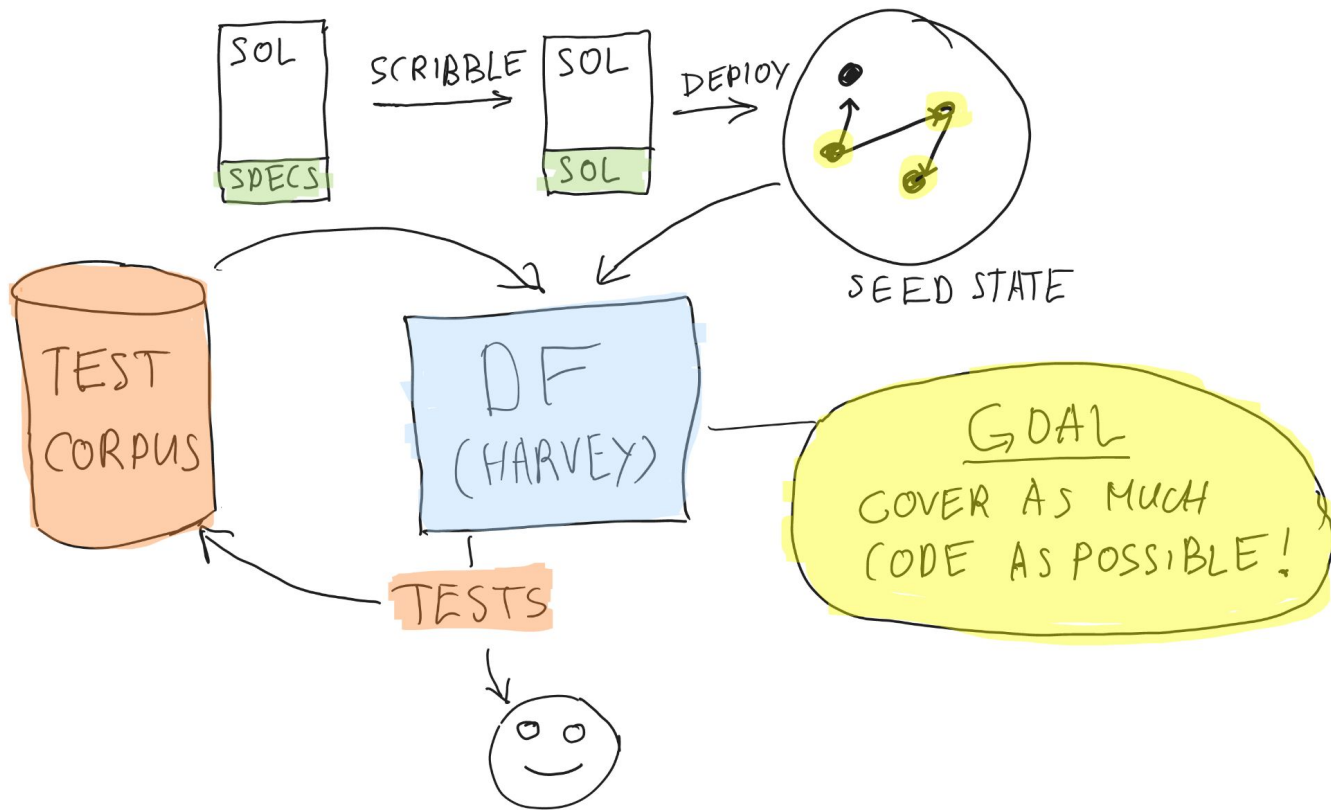
# Fuzzing contract systems

Typical workflows





# Incremental fuzzing



# Incremental fuzzing

- Enables **iterative changes** to code and specifications
- Provides **quick feedback**
- Reuses **existing corpus**

# Typical workflow for new codebase

- 1) Set up fuzzer
- 2) Start fuzzing campaign
- 3) Write or refine properties (if necessary modify code)
- 4) Stop fuzzing campaign and review coverage, etc.
- 5) If necessary refine setup, goto (2)

# How to increase coverage using fuzzing lessons

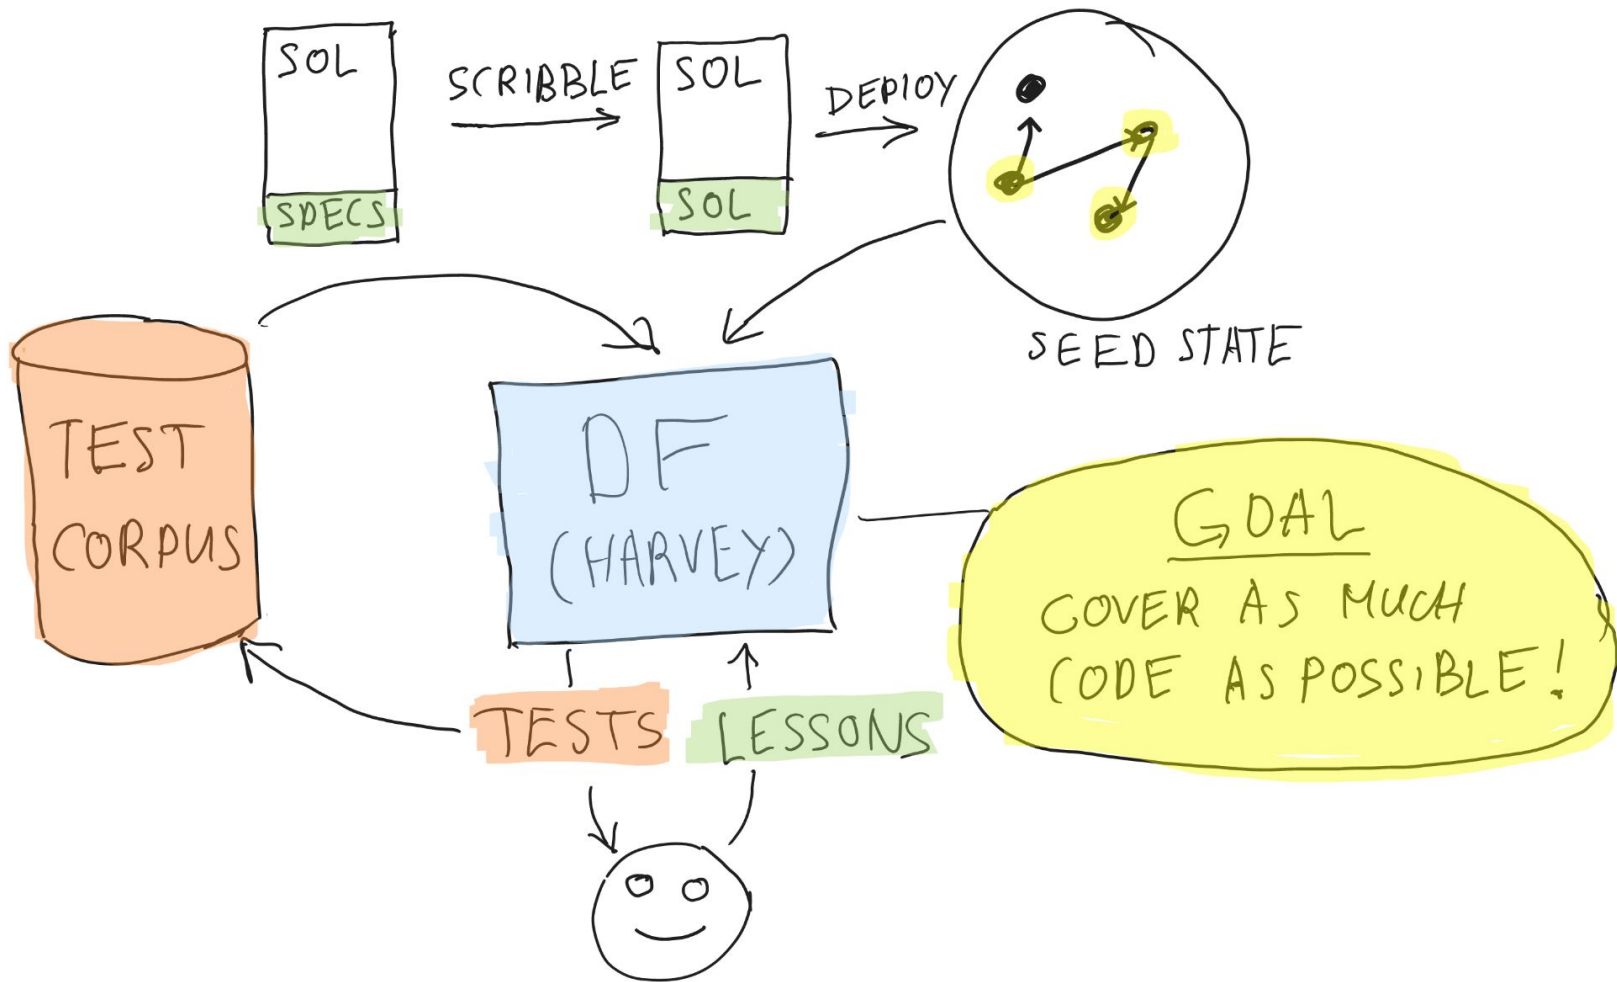
Helping the fuzzer one lesson at a time

# Problem: some code isn't covered

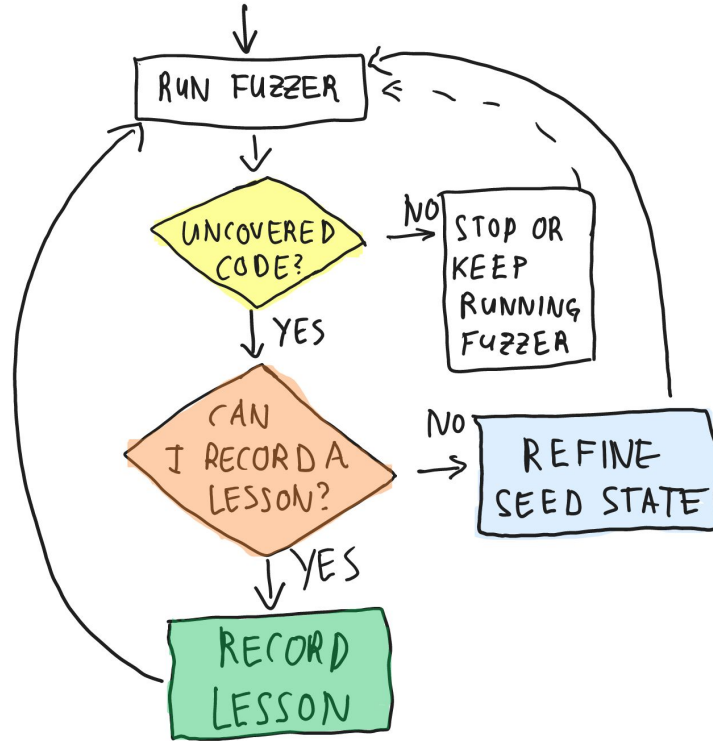
1. Make sure you have let the fuzzer run for some time
2. Review your fuzzer setup:
  1. What transactions are necessary to cover the code?
  2. Are all invoked contracts declared as “contracts under test”?
  3. Are preconditions for the transactions met? Missing token approvals?
  4. ...

# Fuzzing lessons: key idea

- Manually **suggest a sequence of transactions**
- **Concretely:** Write a script and let the fuzzer **observe the executed transactions**
- **Restart the fuzzer** with additional suggestions



# Workflow for increasing coverage





# Example

contracts/contracts/lesson.js  
Main 22 lines

```
1  pragma solidity ^0.4.18;
2
3  contract GasLimit {
4      bool isDestroyable;
5      address owner;
6      constructor(address _owner) public {
7          owner = _owner;
8      }
9      function destroy() public {
10         require(!isDestroyable);
11         selfdestruct(owner);
12     }
13     function permitDestroy(uint8 v, bytes32 r, bytes32 s) external payable {
14         require(!isDestroyable);
15         bytes32 hash = keccak256(abi.encode("permit-destroy", address(this), block.chainid));
16         address signer = ecrecover(hash, v, r, s);
17         require(signer != address(0x0));
18         require(signer == owner);
19         isDestroyable = true;
20     }
21 }
22
23
```

```
$ fuzz lesson start --description "..."
```

```
$ npx hardhat run --network localhost scripts/lesson.js
```

as owner:

1) hash = keccak256(...)

2) v, r, s = ownerKey.sign(hash)

3) contract.permitDestroy(v, r, s)

```
$ fuzz lesson stop
```

# Example (after recording the lesson)

Contract: GaslessDestroy

Main 22 lines 0 issues Collapse

```
1  pragma solidity 0.8.17;
2
3  contract GaslessDestroy {
4      bool isDestroyable;
5      address owner;
6      constructor (address o) {
7          owner = o;
8      }
9      function destroy() external {
10         require(isDestroyable);
11         selfdestruct(payable(msg.sender));
12     }
13     function permitDestroy(uint8 v, bytes32 r, bytes32 s) external payable {
14         require(!isDestroyable);
15         bytes32 hash = keccak256(abi.encode("permit-destroy", address(this), block.chainid));
16         address signer = ecrecover(hash, v, r, s);
17         require(signer != address(0x0));
18         require(signer == owner);
19         isDestroyable = true;
20     }
21 }
22
23
```

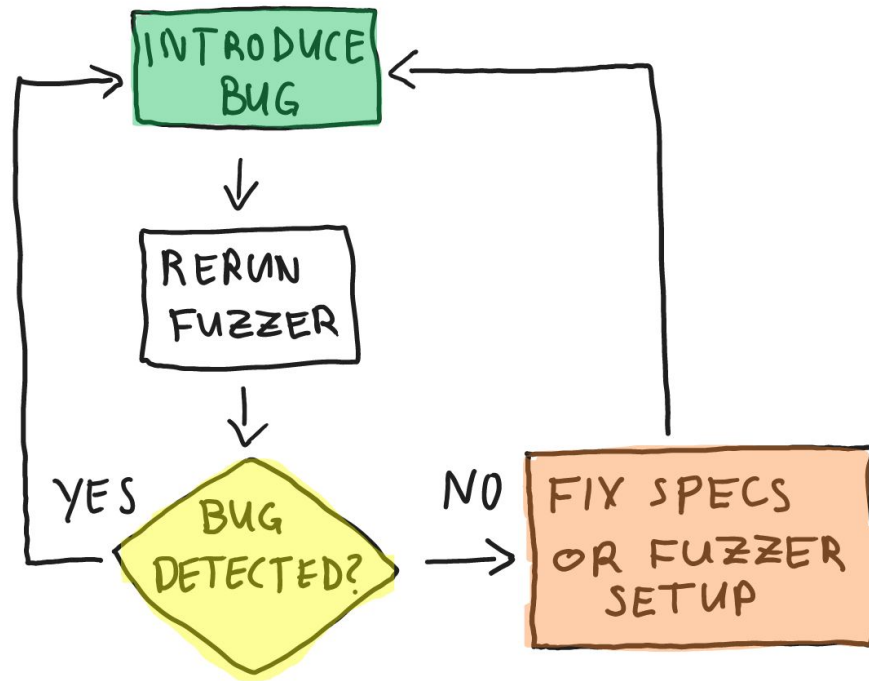
# Testing specifications and fuzzing setup

A good specification catches bugs

# Problem: how good are my specs?

- Related: **How good is my fuzzer setup?**
- Getting specifications “right” is challenging:
  - Code has corner cases, and so do specifications
  - **You can always write more specifications**
  - Tradeoff: **precision vs. effort**

# Workflow for testing specifications



# Example

```
/// #if_succeeds old(balanceOf(t)) <= balanceOf(t);  
function transfer(address t, uint256 a) public {  
    address owner = _msgSender();  
    _transfer(owner, t, a);  
    return true;  
}
```



No bug!

# Example: bug 1

```
/// #if_succeeds old(balanceOf(t)) <= balanceOf(t);  
function transfer(address t, uint256 a) public {  
    address owner = _msgSender();  
    _transfer(owner, t, 1);  
    return true;  
}
```



No bug!

# Example: refined specification 1

```
/// #if_succeeds balanceOf(t) == old(balanceOf(t)) + a;  
function transfer(address t, uint256 a) public {  
    address owner = _msgSender();  
    _transfer(owner, t, a);  
    return true;  
}
```



Bug!



# Example: refined specification 2

```
/// #if_succeeds t == msg.sender ||  
/// balanceOf(t) == old(balanceOf(t)) + a;  
function transfer(address t, uint256 a) public {  
    address owner = _msgSender();  
    _transfer(owner, t, a);  
    return true;  
}
```



No bug!

# Fuzzing under the hood

Challenges when fuzzing smart contracts

# Harvey

- Fuzzer for the **Ethereum Virtual Machine (EVM)** under development since **September 2017**
- Acquired by **ConsenSys** in 2018
- Integrated in **MythX** and **Diligence Fuzzing** products
- Regularly used in **audits** and **client engagements**
- Integrates **novel fuzzing techniques** developed in collaboration with Maria Christakis from MPI-SWS

# What's greybox fuzzing?

- Between **blackbox** (without feedback) and **whitebox** fuzzing (a.k.a. symbolic execution)
- Runs **concrete** inputs instead of **symbolic** inputs
- Uses **lightweight feedback** to guide fuzzer
- **No constraint solving**

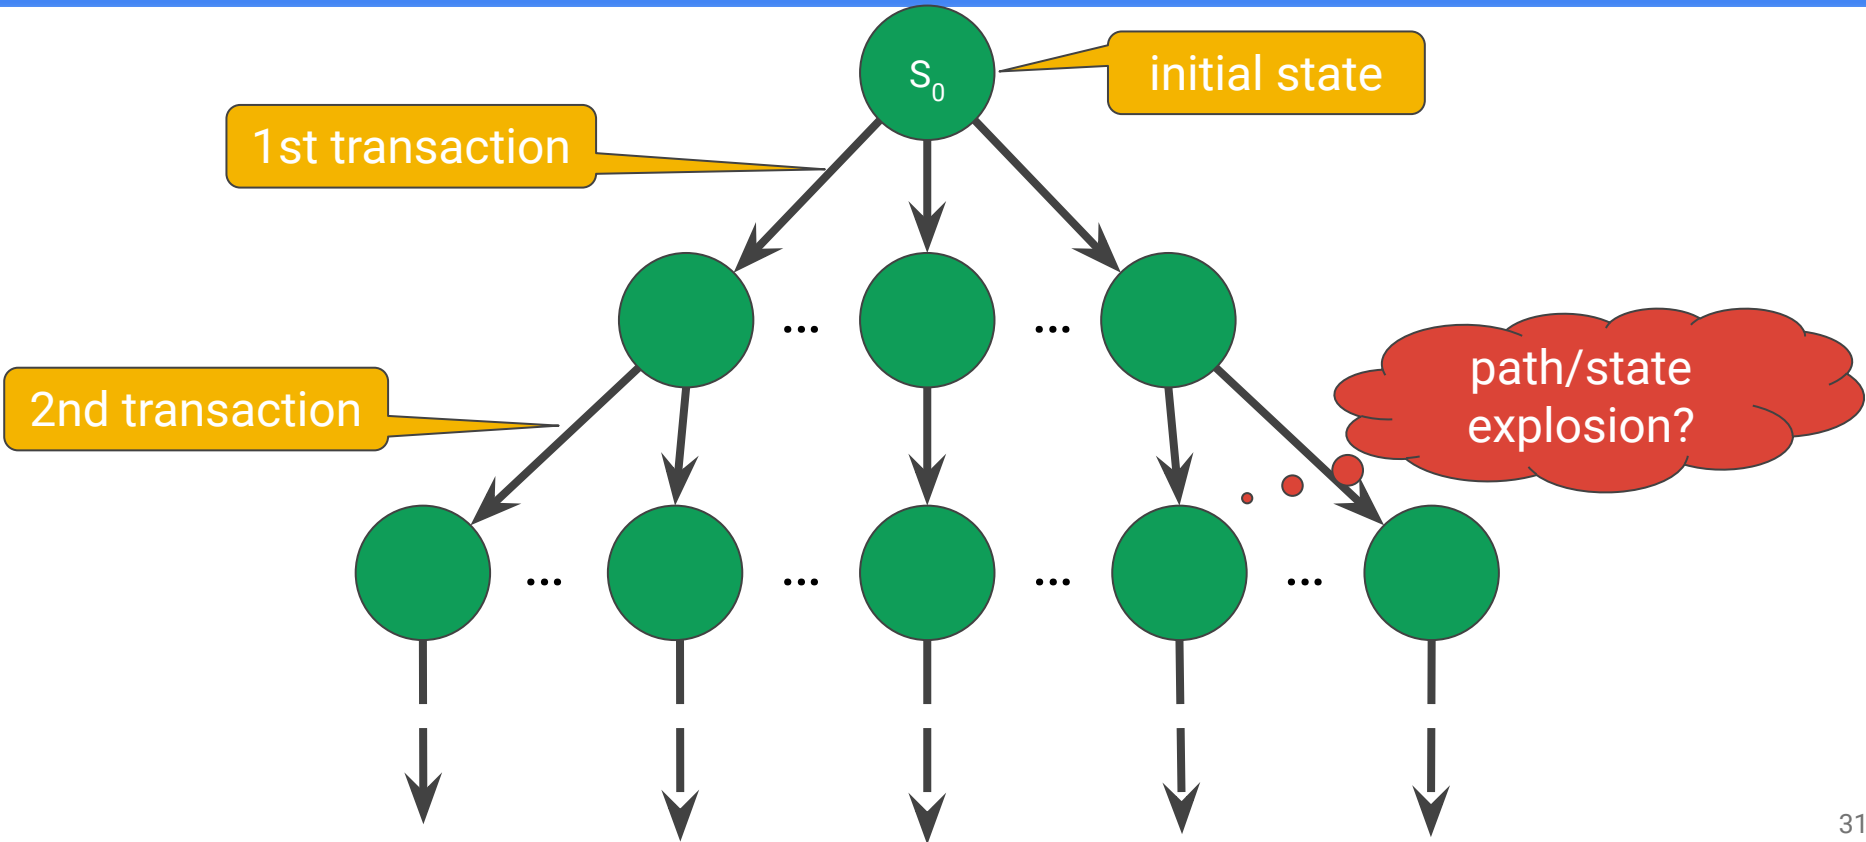
# Greybox fuzzing algorithm

```
queue := run_seed_input(program, seed)  
while (not interrupted) {  
    i := select_input(queue); e := assign_energy(i)  
    while ( $0 < e$ ) {  
        f := fuzz_input(i); path_id := run_input(program, f)  
        if path_id not in queue { queue[path_id] := f }  
        e := e - 1  
    }  
}  
return queue
```

# High-level idea for fuzzing contracts

- User provides **initial state** (incl. one or more contracts)
- User provides custom **correctness properties** (optional)
- Fuzzer **generates sequences of transactions** invoking contracts under test:
  - mutates transaction data (e.g., function inputs)
  - mutates sequences

# Execution model



# Fuzzing sounds simple, but ...

- ... **efficient fuzzing (of contracts) is challenging!**
- Harvey incorporates several **novel fuzzing technique**
- Long-term research effort:

Valentin Wüstholtz and Maria Christakis. **Harvey: A Greybox Fuzzer for Smart Contracts** (ESEC/FSE 2020)

Valentin Wüstholtz and Maria Christakis. **Targeted Greybox Fuzzing with Static Lookahead Analysis** (ICSE 2020)

...



# Challenge 1: narrow checks

```
function Bar(int256 a, int256 b, int256 c) returns (int256) {  
    int256 d = b + c;  
    if (d < 1) {  
        if (b < 3) { return 1; }  
        if (a == 42) { return 2; }  
        return 3;  
    } else {  
        if (c < 42) { return 4; }  
        return 5;  
    }  
}
```

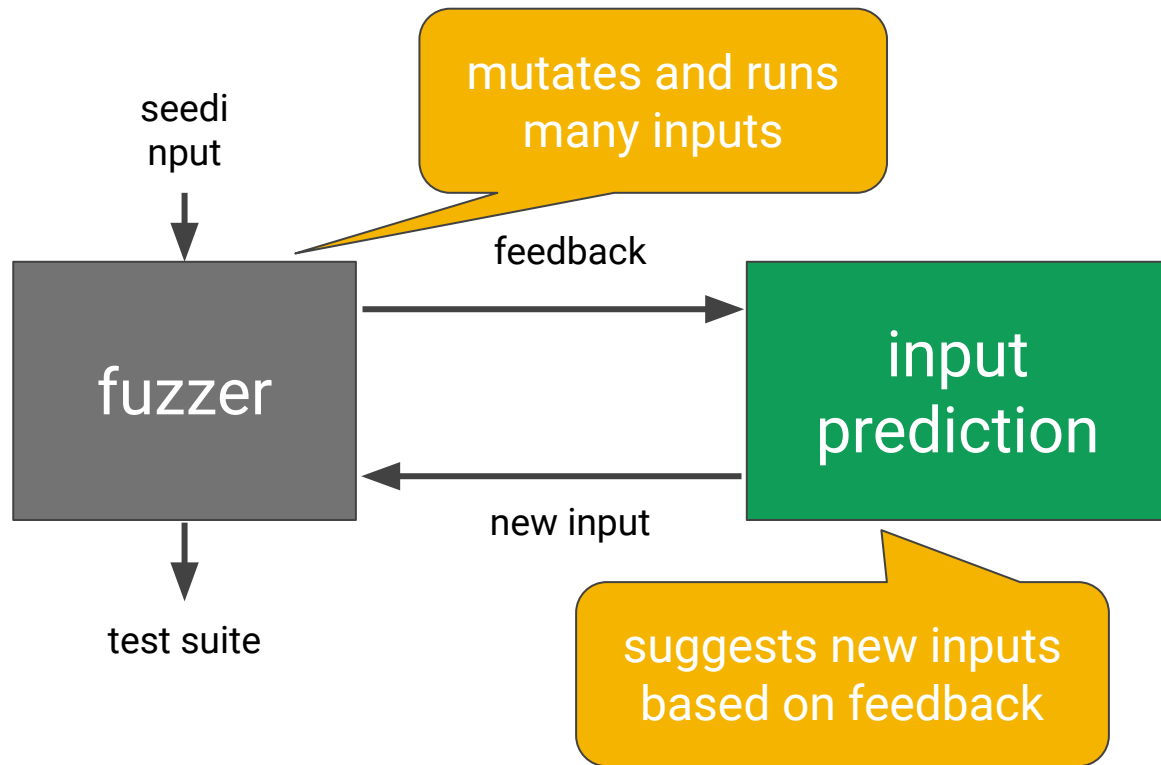
almost impossible to reach  
by random mutations

addressed by  
input prediction  
technique

# Input prediction: key idea

- Make greybox fuzzing “more white” by **collecting additional feedback**
- **Approach:** use feedback to suggest new inputs
- **Feedback:** how far is the execution from “flipping” each branch (**branch distance**)

# Fuzzing + Input Prediction = Coverage!



# Branch distance: example

```
function Bar(int256 a, int256 b, int256 c) returns (int256) {  
    int256 d = b + c;  
    if (d < 1) {  
        if (b < 3) { return 1; }  
        if (a == 42) { return 2; }  
        return 3;  
    } else {  
        if (c < 42) { return 4; }  
        return 5;  
    }  
}
```

distance for a==0, b==3, c== -3?

42!

# Branch distance: example after fuzzing

```
function Bar(int256 a, int256 b, int256 c) returns (int256) {
```

```
    int256 d = b + c;
```

```
    if (d < 1) {
```

```
        if (b < 3) { return 1; }
```

```
        if (a == 42) { return 2; }
```

```
        return 3;
```

```
    } else {
```

```
        if (c < 42) { return 4; }
```

```
        return 5;
```

```
    }
```

```
}
```

distance for a==7, b==3, c== -3?


35!

# What can we learn from the feedback?

- Input  $a$  **taints the branch**
- What input  $a'$  should we try next?
  - What value  $a'$  makes distance function  $f(a') == 0$ ?
  - Sounds familiar? Essentially **root-finding**!

# Secant Method

- Iteratively finds roots
- **No derivatives needed!**



# What if distance function isn't linear?

- Apply **iteratively!**
- Works for non-linear conditions

Example:  $x^{**4} + x^{**2} == \dots$

- **Fun fact: one step is often enough**

(average success rate of **99%** in our experiments)



# Effectiveness of input prediction

- Finds **bugs** orders of magnitude faster (**~5x**)
- Increases **coverage** significantly (up to **~3x**)

# Challenge 2: deep vulnerabilities

```
contract Foo {
```

```
  int256 private x;
```

```
  int256 private y;
```

```
  function Bar() public { assert(x != 42); }
```

```
  function SetY(int256 ny) public { y = ny; }
```

```
  function IncX() public { x++; }
```

```
  function CopyY() public { x = y; }
```

```
}
```

persistent state variables

requires multiple calls to trigger


addressed by  
demand-driven  
fuzzing

# Demand-driven fuzzing: key idea

- Fuzzing long transaction sequences is expensive:
  - More inputs to fuzz
  - Longer execution
  - Number of paths grows exponentially
- **Idea:** avoid it unless it may improve coverage

# Coverage of what?

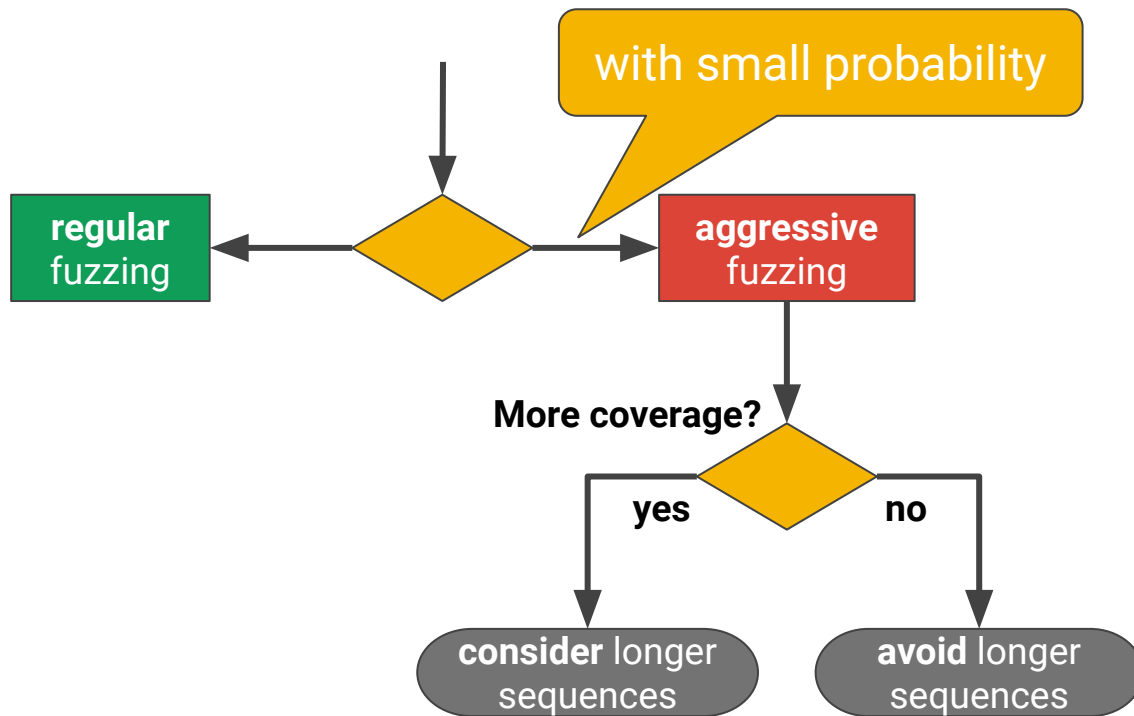
- **Issue:** increasing coverage is trivial if paths span multiple transactions
- Harvey only considers **path of last transaction**



# Aggressive fuzzing

- Use aggressive fuzzing to determine if coverage may increase and longer sequences should be considered
- Main difference: **allow fuzzing the persistent state**

# Regular vs. aggressive fuzzing



# Effectiveness of demand-driven fuzzing

- ~**60%** of bugs require multiple transactions
- Finds **bugs** orders of magnitude faster (~**3x**)
- Increases **coverage** significantly

# Summary

*Shared some tips and tricks for fuzzing smart contract systems effectively*

